

Problem Set 2: Part I

Problem 1: Fixed-length and variable-length records

1.1 and 1.2

record contents

| | | | |
|-------|--------------|------|-------|
| 15172 | Barbie#----- | 2023 | PG-13 |
|-------|--------------|------|-------|

length in bytes

| |
|----|
| 34 |
|----|

show how you computed the length:

| |
|--|
| For fixed-length records, the stored values must take up the amount of space specified in type definition. So, the length in byte = $5 + 20 + 4$ (4 bytes of int) + $5 = 34$. |
|--|

1.3 and 1.4

record contents

| | | | | | | | |
|-------|---|--------|---|------|---|-------|---|
| 15172 | # | Barbie | # | 2023 | # | PG-13 | # |
|-------|---|--------|---|------|---|-------|---|

length in bytes

| |
|----|
| 28 |
|----|

show how you computed the length:

| |
|--|
| For variable-length records, I choose to terminate field values with a special delimiter character '#', each of which is 2 bytes. Therefore, the length in byte = $5 + 2 + 6 + 2 + 4 + 2 + 5 + 2 = 28$. |
|--|

1.5 and 1.6

record contents

| | | | | | | | | |
|----|----|----|----|----|-------|--------|------|-------|
| 10 | 15 | 21 | 25 | 30 | 15172 | Barbie | 2023 | PG-13 |
|----|----|----|----|----|-------|--------|------|-------|

length in bytes

| |
|----|
| 30 |
|----|

show how you computed the length:






| |
|---|
| For variable-length records that begins with a header of offsets, there are (#fields+1) headers, i.e. $4+1 = 5$ headers. Given that each integer metadata takes up 2 bytes, then the value of the first field will start at position 10, the second at position $10+\text{len}('15172')=15$, etc. The last header offset stores the end position (the length) of this record, which is 30. |
|---|

1.7

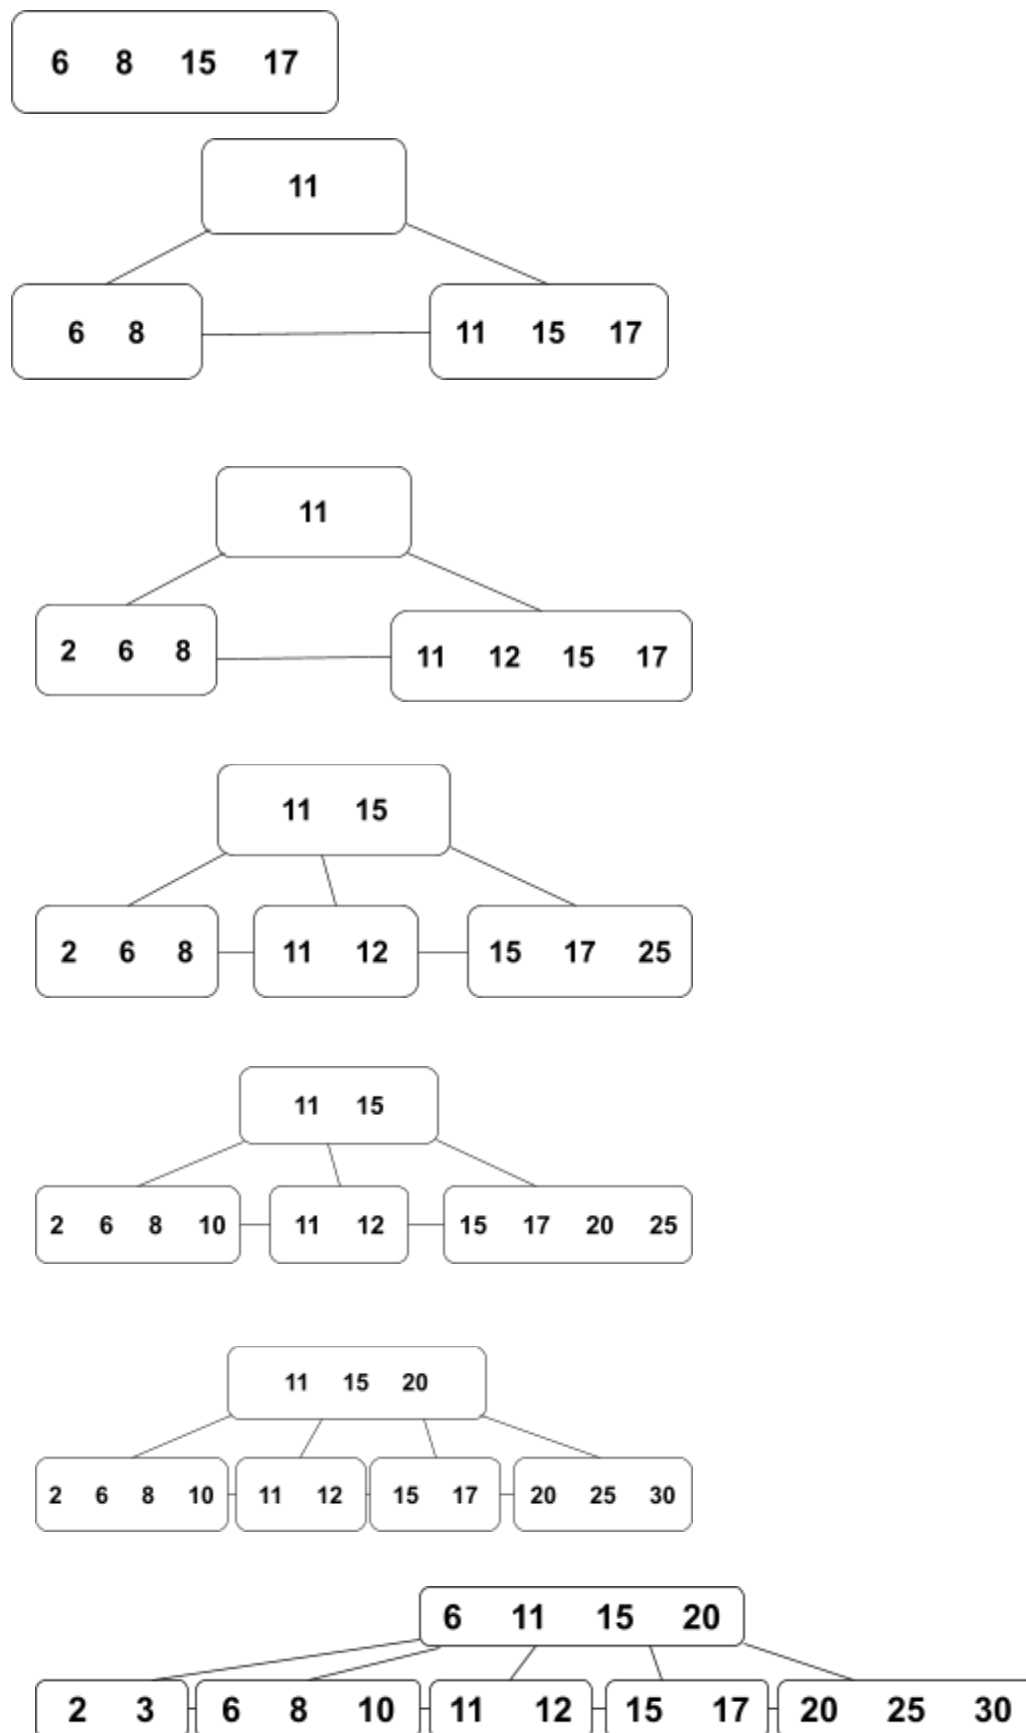
record contents


| | | | | | | | |
|----|----|----|----|----|-------|------------------|-------|
| 10 | 15 | -1 | 31 | 36 | 87654 | The Color Purple | PG-13 |
|----|----|----|----|----|-------|------------------|-------|

Problem 2.1: Insertions into a B-tree (order = 2)



Problem 2.2: Insertions into a B+tree





Problem 2.3: Insertions into a linear hash table

a bucket is added whenever the number of items in the table exceeds three times the number of buckets.

6 = 0110

15 = 1111

8 = 1000

17 = 0001 0001

11 = 1011

12 = 1100

2 = 0010

25 = 0001 1001

20 = 0001 0100

10 = 1010

30 = 0001 1110

3 = 0011

27 = 0001 1011

$i = \text{ceil}(\log_2 n) = \text{ceil}(\log_2(5)) = 3$ (11)

before first increase

| | |
|----|------------|
| 00 | 6, 8, 12 |
| 01 | 15, 17, 11 |

after first increase

| | |
|----|------------|
| 00 | 8, 12 |
| 01 | 15, 17, 11 |
| 10 | 6, 2 |

before second increase

| | |
|----|----------------|
| 00 | 8, 12, 20 |
| 01 | 15, 17, 11, 25 |
| 10 | 6, 2 |

after second increase

| | |
|----|-----------|
| 00 | 8, 12, 20 |
| 01 | 17, 25 |
| 10 | 6, 2, 10 |
| 11 | 15, 11 |

before third increase

| | |
|----|--------------|
| 00 | 8, 12, 20 |
| 01 | 17, 25 |
| 10 | 6, 2, 10, 30 |
| 11 | 15, 11, 3 |

after third increase

| | |
|-----|---------------|
| 00 | 8 |
| 01 | 17, 25 |
| 10 | 6, 2, 10, 30 |
| 11 | 15, 11, 3, 27 |
| 100 | 12, 20 |

final state of the table

| | |
|---|---------------|
| 0 | 8 |
| 1 | 17, 25 |
| 2 | 6, 2, 10, 30 |
| 3 | 15, 11, 3, 27 |
| 4 | 12, 20 |